

An Interactive Visualizer for Equality Solvers using E-Graphs

Princeton University COS 516 Course Project (Fall '22)

Ryan Torok
Princeton University

Leon Schürmann
Princeton University

1 Introduction

The *Theory of Equality* (TOE) is a theory in first-order logic which axiomatizes that the equality operator ($=$) behaves reasonably [2]. A well-known algorithm for finding all equivalence classes of a given formula, which constitutes a theory solver for TOE, is to list all terms that appear in the formula, then iterate through pairs of terms, adding terms to equivalence classes using a union-find data structure by repeatedly applying the rule $a = b \rightarrow f(a) = f(b)$, where a and b are terms, and f is a function symbol. In the worst case, running this procedure until no more matches are possible requires quadratic time. While this is still a polynomial-time algorithm, it is rather naive in its method to choose new pairs and it is definitely possible to do better.

E-graphs [3] improve the time complexity of this procedure by representing the terms of an expression as a directed acyclic graph (DAG), such that constant or function symbols are represented using vertices, and the composition of terms is given through the edges. Each vertex belongs to a given *equality class*, which defines subterms that are equivalent as per equivalence under the TOE. Edges in this graph lead from a given vertex v , parameterized by an argument index $0 \leq i < \text{arity}(v)$ towards an equivalence class, which binds the symbol's n th parameter to any element of this class. For example, Figure 1 represents an E-graph for the formula $(a * 2)/2$.

This structure is an inherently compact representation of possible alternate expressions of a *program* (defined as the top-level term in the provided expression). This characteristic makes it possible to efficiently find all equality classes without requiring the linear pass over all terms. In addition to the equivalences established




Figure 1: E-graph representing the formula $(a * 2)/2$ [4]

by the equivalence operator as axiomatized by the TOE in first-order logic with equality, E-graphs can set additional expressions equal by a user-provided, arbitrary set of additional rewrite rules (unidirectional definitions of allowed term transformations). This makes E-graphs a powerful tool for many applications such as optimizing compilers, where there are a variety of ways to express and equivalent concept. Given that E-graphs can encompass a multitude of equivalent programs, this enables compilers to search them for expressions which are ideal under a set of constraints, for instance in an effort to eliminate common sub-expressions. This particular application area of E-graphs is called *equality saturation*.

egg is a popular Rust library that implements E-graphs, and has seen widespread use in the research community, which can be attributed to its flexibility, efficiency and ergonomic developer interface [4].

Unfortunately, the precise mechanisms employed both within the congruence closure algorithm, as well as the extended e-graph data structures, can be unintuitive for

novice users. We presume that an interactive visualization can help demonstrate the basic concepts behind reasoning about equality under a given set of rewrite rules. In this project, we create *eggviz*, a web-based visualization tool that allows the user to interactively explore the equivalence-matching procedure and program optimization techniques in a step-by-step manner.

2 User Interface Overview

We begin by providing a high-level overview over the feature-set and ideas incorporated into *eggviz*. Its user interface follows a layout inspired by the CDCL solver as introduced in the course¹. Figure 2 contains a screenshot providing a high-level overview of the main *eggviz* user interface components. In its top-bar, a user can enter a program following a LISP-inspired syntax. We choose this syntax as it is easy to parse (including implementing full syntax error handling) and simple for the user to reason about its equivalent E-graph. We define our LISP syntax formally in Section 3.1.

Through the application’s left-hand control pane, the user can add *rewrite rules* to be applied to the program, using the same LISP-like syntax as defined in 3.1. Following adding a program and at least one rewrite rule, the user can press the *Graph!* button to show an E-graph representation of the program. The graph, which we render through the *vis.js* [1] library, represents each equivalence class (E-class) using a different color. Unfortunately, *vis.js* does not support a grouping a set of vertices into a class, as would be particularly useful to illustrate the concept of equivalence classes. While extending *vis.js* to support such a visualization mode is outside of the scope of this project, this could be a particularly useful future task. We work around this limitation by introducing two distinct types of nodes: *box-shaped nodes*, which represent individual terms and *circular nodes*, which serve as a root node for each equivalence class. Along with this distinction, we further establish two types of relations, modeled through edges: terms in an equivalence class are connected by bold edges without arrows to the root node of that class. The thinner, directed edges point from a box-shaped node to a root node of the equivalence class of the subterm(s) of the corresponding expression. These types of edges are also marked with a number to

correspond to the argument index of that subterm within the parent term.

eggviz supports two usage modes: first, clicking on a rewrite rule at the left pane will apply that rewrite rule iteratively until it can no longer be applied. For this feature, it is up to the user to decide in what order to apply the rewrite rules to generate the desired equivalences. The second mode of user interaction is automatically applying all rewrite rules once, by clicking the *Auto* button. This has the same effect as clicking all of the rewrite rules in order. In other words, for a program with two rewrite rules R_1 and R_2 , *eggviz* would apply R_1 everywhere it can, then do the same with R_2 , but it will not return to R_1 afterwards, even if R_2 created new instances where it could be applied (unless, of course, the user clicks *Auto* or R_1 again).

3 Implementation and Internals

eggviz is powered internally by *egg*, a Rust E-graph library [4]. To integrate *egg* with *eggviz*, we implement a parser for our LISP-like syntax, a dedicated *single-step scheduler* that allows the user to advance the state of the E-graph in a step-wise manner, and a formatting mechanism for translating the graph state to JavaScript in a Rust library which we import into our web page as a WebAssembly module. We also implement and set of JavaScript functions responsible for taking the graph state returned by Rust and drawing the graph for the user using *vis.js*. In this section, we perform a deep dive on the design and implementation of each of these features.

3.1 LispyLang and Parsing

LispyLang is the name we give to the LISP-like syntax we introduced in the previous section and that *eggviz* uses for its program and rewrite rules. LispyLang implements the following syntax:

Program : Term
Term : Symbol | (Symbol Term ...)

A *Symbol* can be any string of characters that does not contain either of the characters "(", " ", or ")". In other words, all variables, numbers, and similar are all considered symbols, since neither *eggviz* nor *egg* assume any particular arithmetic, so the symbol 1 has no

¹<https://www.cs.princeton.edu/courses/archive/fall22/cos516/cdcl/>

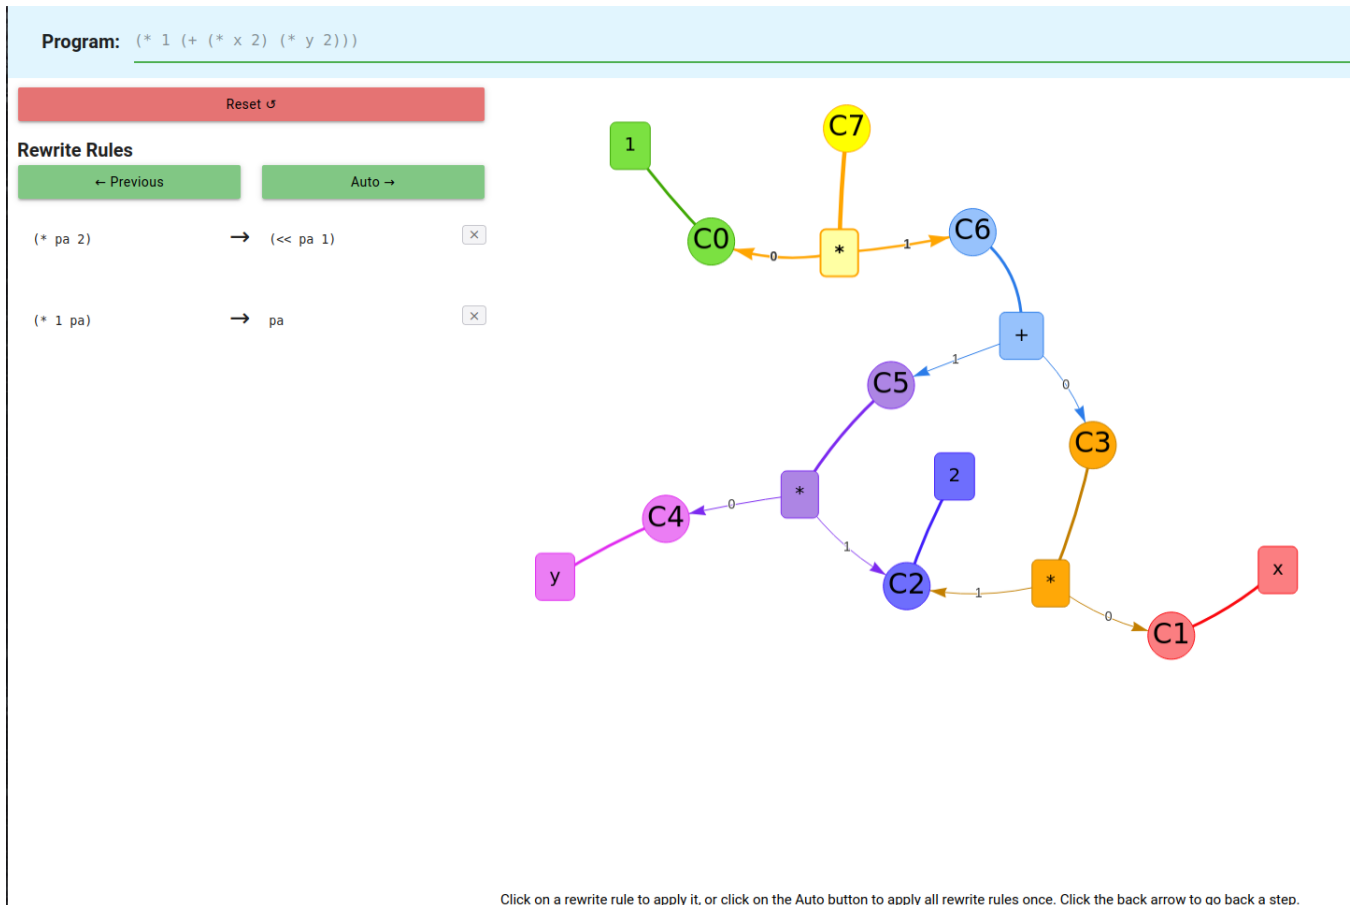


Figure 2: The `eggviz` interface. The program is entered at the top, the rewrite rules at the left, and the graph is shown in the main display area.

more inherent meaning than the symbol x . Likewise, the behavior of function symbols is determined entirely through rewrite rules, meaning the operator $+$ could be defined to just produce the value 0 always or even take three arguments.

However, there are two rules `eggviz` does enforce with regard to symbols: first, all applications of a symbol occurring across the program and all rewrite rules must have the same *arity*, or number of arguments. For example, if a program included the term $(+ a b)$, then the terms $(+ 1 2 3)$, $(+ 1)$, or $+$ could not be used elsewhere in the program or in any rewrite rule, since the $+$ operator was already declared with an arity of 2.

Second, there is one type of symbol that is assigned special meaning: *generic variables*. By convention, these are any symbol beginning with the character `p`. Generic variables are used to dictate that a rewrite rule applies for all values of the generic variable. For example, the identity property of the multiplication operator $*$ can be written as the rewrite rule $(* 1 pa) \rightarrow pa$. This means that the symbol $*$ applied to the symbol 1 and any symbol can be rewritten as just that symbol. The definition of generic variables is local to the rewrite rule in which they appear. Likewise, to avoid confusion, variables beginning the character `p` are not permitted within the program expression.

We designed `LispyLang` with a LISP-like syntax with prefix notation to aid parsing and to make clear that no operator has any inherent meaning in the system, including $*$ and $+$. We implement a simple two-function top-down parser for `LispyLang` that supports full handling of syntax errors (which are written to the footer of the UI if they occur).

3.2 Rewrites and the Single-Step Scheduler

When the user indicates they want to graph a program, once the program and rewrite rules are parsed and validated for syntax and pass the arity checker, `eggviz` takes its internal representation of the program and rewrite rules and produces a string with the syntax `egg` expects, which differs from `LispyLang` only in its notation for generic variables. Next, `eggviz` needs to respond to which rewrite rules the user chooses to apply.

Unfortunately, `egg` was not designed to be able to selectively apply a single rewrite rule at a time; it only exposes an API to repeatedly apply all rewrite rules until it can no longer find any new equality relations. At

the start of the project, we believed we would need to modify the `egg` library in order to implement the interactivity we wanted for `eggviz`. However, it turned out to be possible to allow selectively applying rewrite rules using the unmodified `egg` library by using (abusing?) one feature of `egg`'s API, the *scheduler*.

The scheduler is a feature of `egg` intended to provide users control over the precise equality saturation strategy used in `egg`'s iterations over the E-graph. This is particularly useful given that there are combinations of programs and rewrite rules for which, depending on the heuristic used to determine which rewrites count as *simplifying*, the simplification process would never finish. For example, a bad simplification heuristic would allow recursive rewrite rules like $pa \rightarrow (+ pa pa)$ to be applied forever. While `egg` can be provided *limits* (such as E-graph size and saturation time) at which it will terminate the equality saturation process, the scheduler gives developers control of when which rewrite rules will be applied and can, for instance, aid in early termination of the equality search after a certain point in these runaway cases.

`egg` further allows passing a custom scheduler by implementing the `Scheduler` trait, which we leverage for `eggviz` by implementing the `SingleStepScheduler` type. Rather than just being intended to prevent runaway rule applications, the single-step scheduler is designed to selectively and iteratively apply rewrite rules based on the user's interaction with `eggviz`. Specifically, it is programmed to take a rewrite rule and apply only that rule for a single iteration, or, if the user clicks the Auto button, apply all rules for a single iteration.

To achieve this, the single step scheduler, on the first iteration, tells `egg` to apply only the rule(s) it wants to. For subsequent iterations, the scheduler tells `egg` not to apply any rule. Then, because `egg` decides to stop when the equality classes are unchanged by an iteration, the process is halted after a single meaningful iteration followed by a no-op. Our design of the single-step scheduler is the key piece enabling us to transform what `egg` implements as an opaque one-shot process into a step-wise procedure that grants the user a great deal of control and ability to experiment with different rewrite rules. We also greatly generified the scheduler and rewrite logic so that these components of the `eggviz` library could be applied by future developers to a different language implementation than `LispyLang` if desired.

3.3 Wasm-JS Integration and Graphing

Using the single-step scheduler, `eggviz` instructs `egg` to apply exactly the rewrite rule(s) the user selects. This allows us to read the state of the E-graph from `egg` after each transformation. This graph state is transformed to a JavaScript object using a series of Maps that encode the relationship between symbols, terms, and expressions. Once we return the graph state to JavaScript, we iterate over this map at multiple levels to draw the graph.

Since E-graphs inherently store the function-argument relationship as an edge between the function symbol and the e-class of the argument (rather than the argument term), we designed our graph to have two types of nodes: one to serve as the edge sink for each e-class, and the other to represent the terms in that e-class. This separation of duty prevents a term whose argument is in an e-class with a large number of members from needing to draw an edge to every member of that e-class, making the graph very hard to read and understand.

One consequence of this design choice is that our displayed graph now has two type of edges, one to represent arguments to a function symbol, and another to connect terms within an e-class. We visually distinguish these two types of edges by making the latter type much thicker and lack arrows. E-classes are also visually represented by the color of their nodes. While we briefly considered a basic *color wheel* algorithm to produce contrasting colors, it was eventually scrapped after we discovered that the `vis.js` library has a built-in coloring feature that is much better at avoiding dark colors that would disguise the text within the nodes.

4 Example

In this section, we provide a short example of how `eggviz` functions. Figure 3 shows `eggviz`'s rendering of the program $(*\ 1\ (+\ (*\ x\ 2)\ (*\ y\ 2)))$, with the rewrite rules $(*\ pa\ 2) \rightarrow (\ll\ pa\ 1)$ and $(*\ 1\ pa) \rightarrow pa$. These rewrite intuitively correspond to the properties "1 times any number is itself" and "any number times 2 is the same as left-shifting that number by one bit", respectively.

Suppose we click on the second rewrite rule to apply it. This tells `eggviz` to find any new equality relations based on the rule that "1 times any number is itself". In this case, this means that the entire program, $(*\ 1\ (+\ (*\ x\ 2)\ (*\ y\ 2)))$, is equivalent to just the term $(+$

$\ (*\ x\ 2)\ (*\ y\ 2))$. In the E-graph, this corresponds to annotating the term rooted at the outermost $*$ operator with the same equality class as the term rooted at the $+$ operator. Figure 4 shows the change to the graph that occurs after applying, this rewrite rule; namely, the outermost $*$ operator, originally part of its own equality class `C7`, was merged into the equality class `C6`, which contains the $+$ term.

Now, suppose we apply the first rewrite rule next. Figure 5 shows the graph after applying this transformation. Two new terms are introduced for $(\ll\ x\ 1)$ and $(y\ \ll\ 1)$, which are included in the same equality classes as the terms $(*\ x\ 2)$ and $(*\ y\ 2)$, respectively. Also, observe that the two \ll nodes are the only new nodes introduced to the graph at this step, since they could use the existing x and y terms, respectively, and the symbol 1 from the outermost $*$ operator as their children.

5 Implementation Breakdown

Both of us contributed equally to this project. Here is a rough breakdown of who worked on what:

- Leon: Rewrite logic, single-step scheduler, language traits, WebAssembly API design, significant refactoring
- Ryan: UI controls and event handlers, `LispyLang` parser and error handling, `egg` library integration, graph beautification
- Pair-programmed: WebAssembly to JavaScript encoding conversion, `vis.js` library integration and graph display

Our implementation of `eggviz` is available on GitHub², with a prebuilt version deployed on GitHub pages³. `eggviz` is licensed under a AGPL-3.0 license and can be freely used, shared, modified and inspected according to the terms of this license.

References

- [1] `vis.js`. <https://visjs.org/>.
- [2] Aaron R. Bradley and Zohar Manna. *The Calculus of Computation*. Springer, New York, 2007.

²<https://github.com/lshuermann/eggviz>

³<https://lshuermann.github.io/eggviz/>

- [3] Greg Nelson. Techniques for program verification. 1981.
- [4] Max Willsey, Chandrakana Nandi, Yisu Remy Wang, Oliver Flatt, Zachary Tatlock, and Pavel Panchekha. egg: Fast and extensible equality saturation. *Proceedings of the ACM on Programming Languages*, 5(POPL):1–29, 01 2021.




Figure 3: eggviz representation of the E-graph for the program $(* 1 (+ (* x 2) (* y 2)))$.




Figure 4: Graph from Figure 3 after applying the rewrite rule $(* 1 pa) \rightarrow pa$.




Figure 5: Graph from Figure 4 after applying the rewrite rule $(\ast \text{ pa } 2) \rightarrow (\ll \text{ pa } 1)$.